

## **Safe harbor statement**

The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, timing, and pricing of any features or functionality described for Oracle's products may change and remains at the sole discretion of Oracle Corporation.

**ORACLE**  
Data Deep Dive  
at AI World

# RAG in a Box

Simplify and Accelerate Data Searches

---

**Massimo Castelli**

*Vice President AI & Data Platform Strategy*

**Doug Hood**

*Product Manager for AI Vector Search and Private AI Services Container*

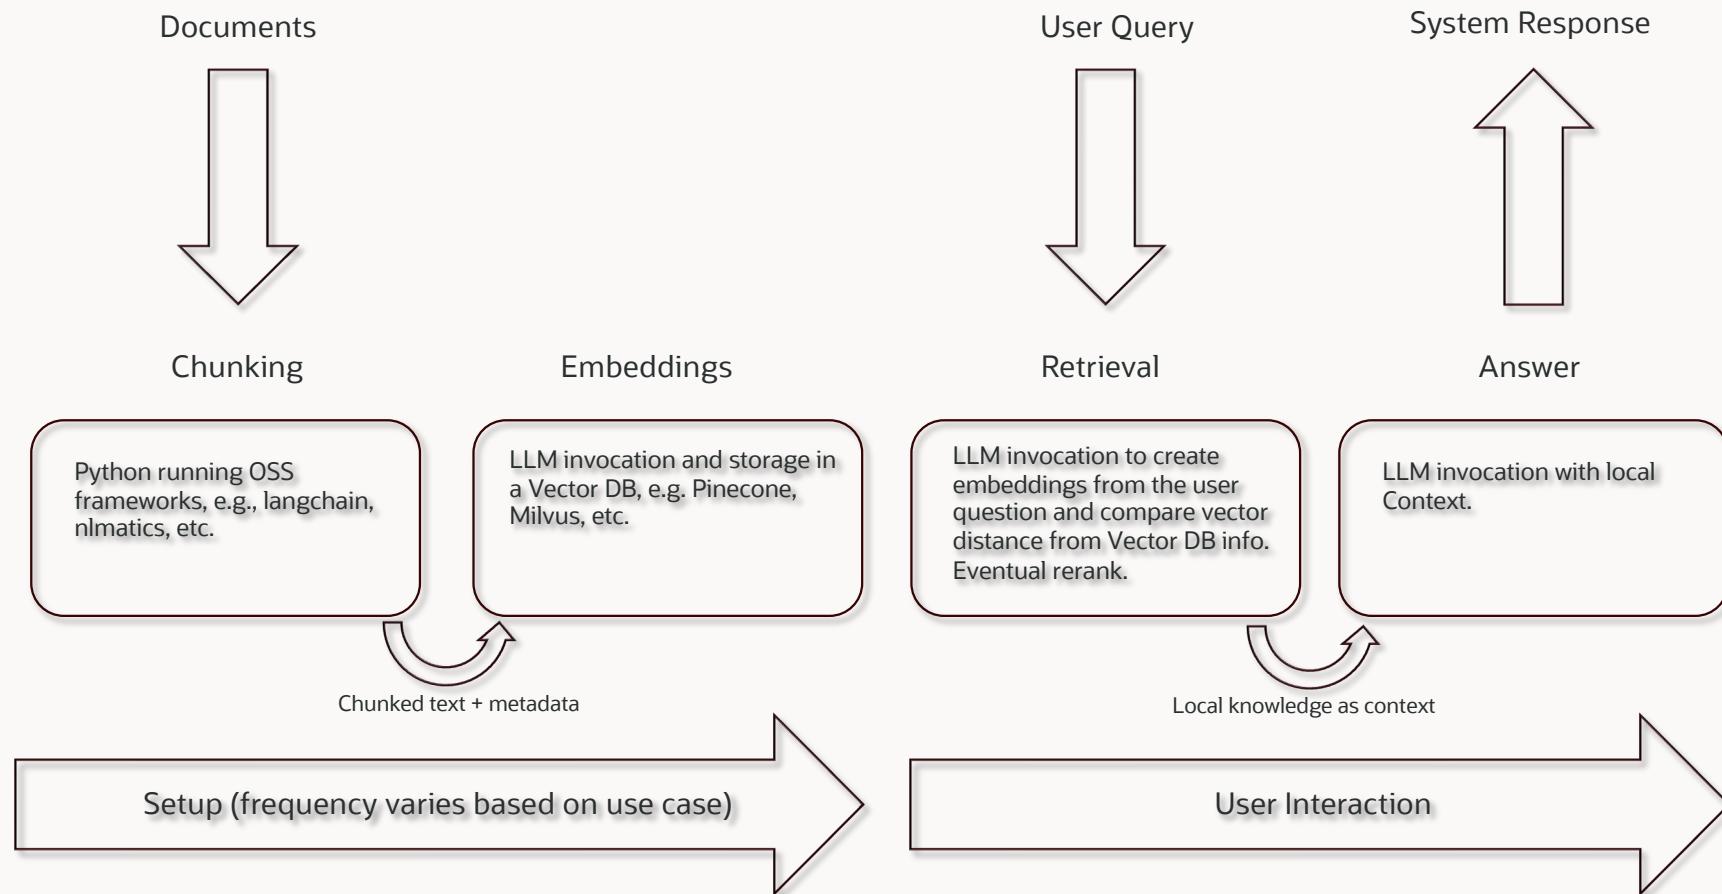


## Retrieval Augmented Generation (RAG)

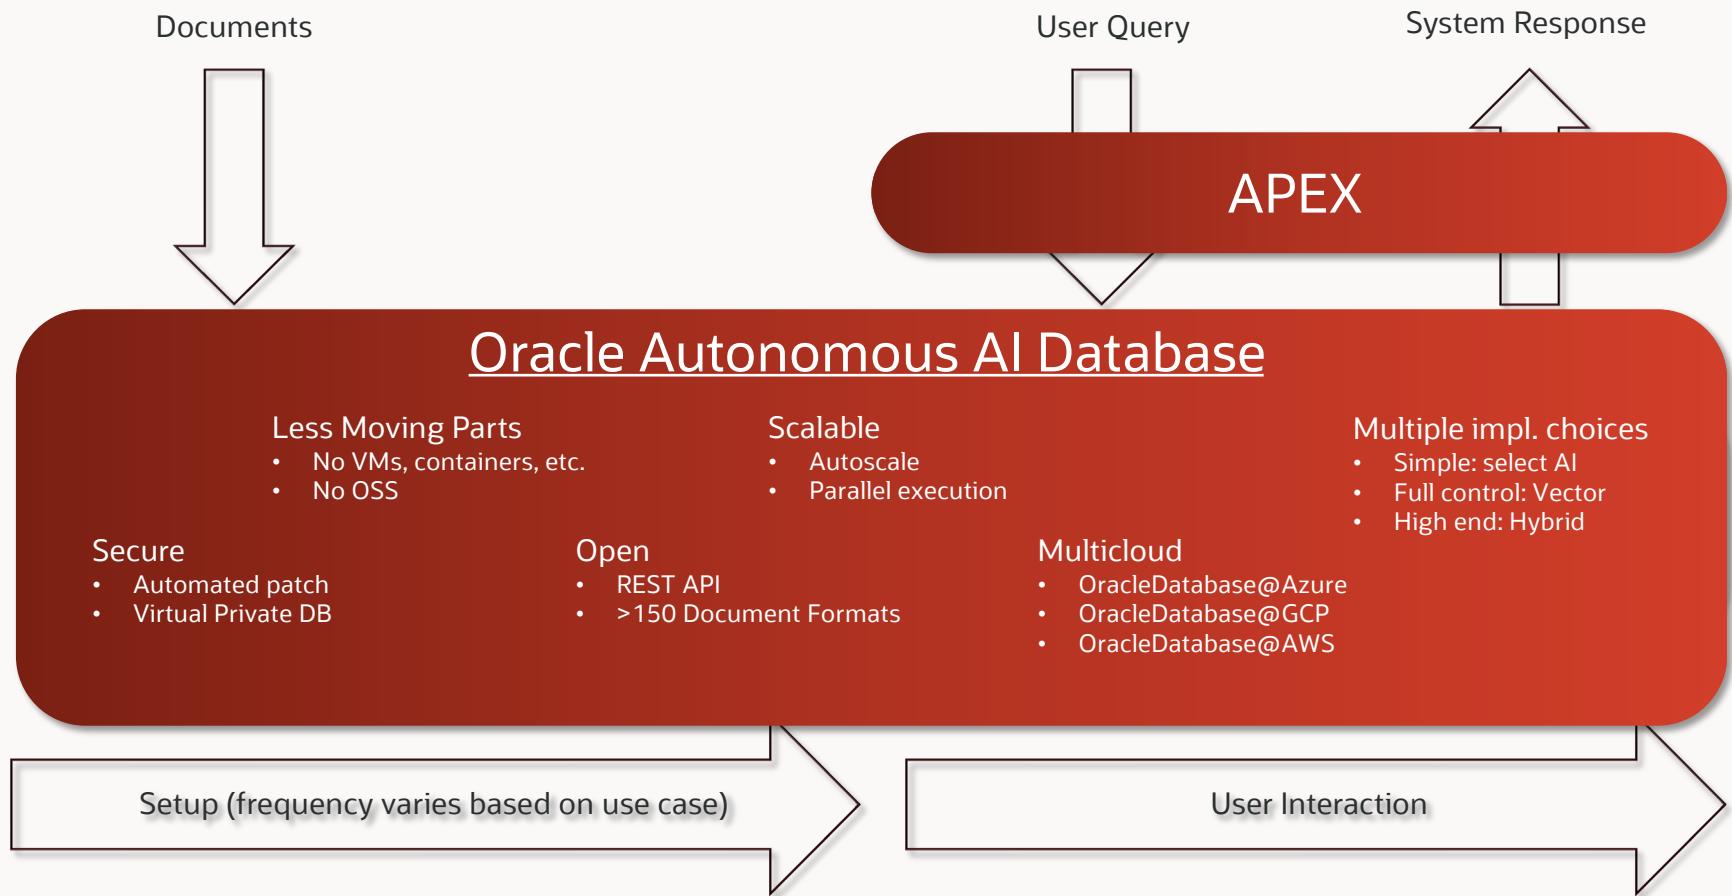
Enables LLMs to use business data to produce better and more contextually relevant answers to user questions while keeping business data secure



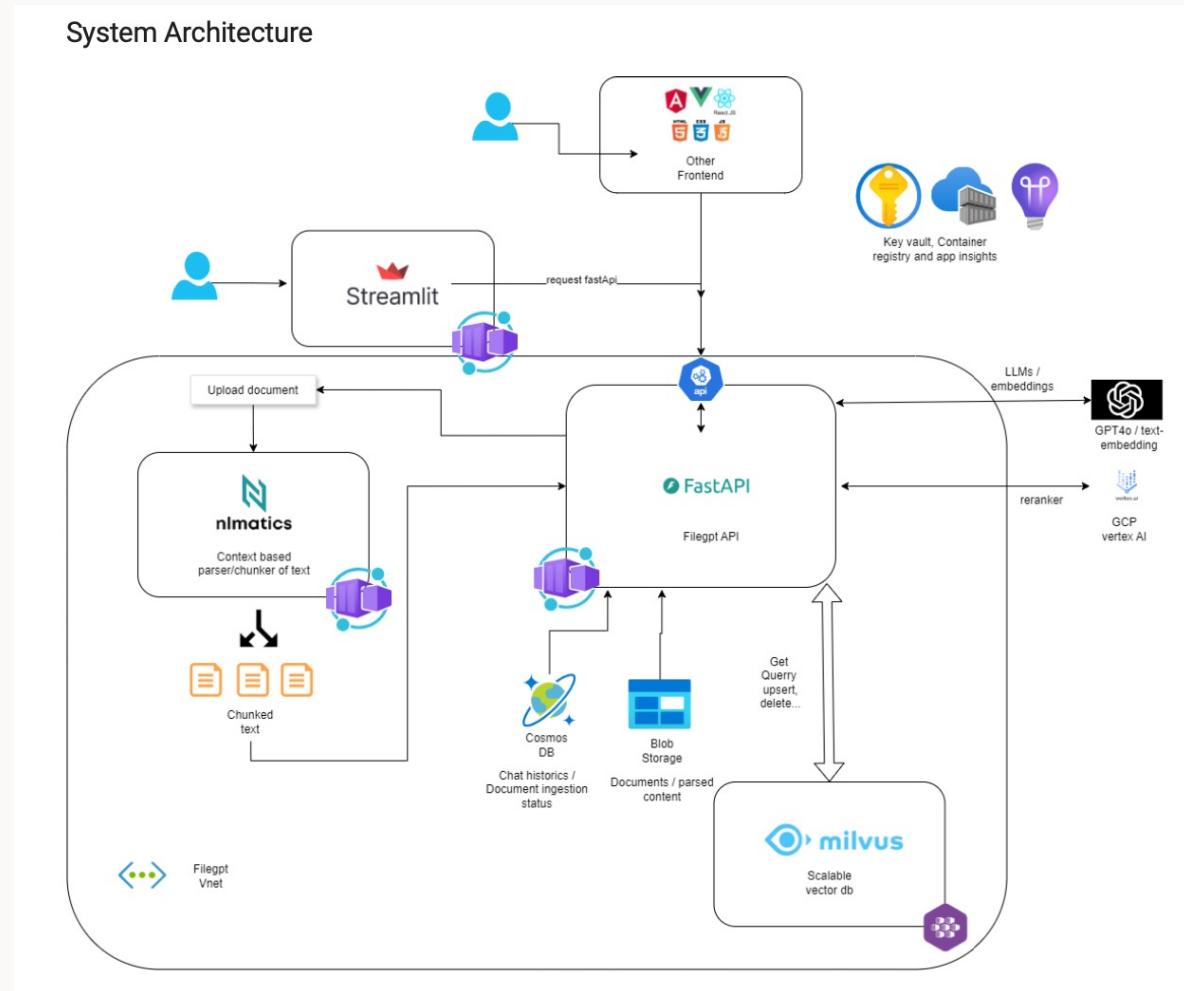
# Typical RAG Pipeline



# Typical RAG Pipeline

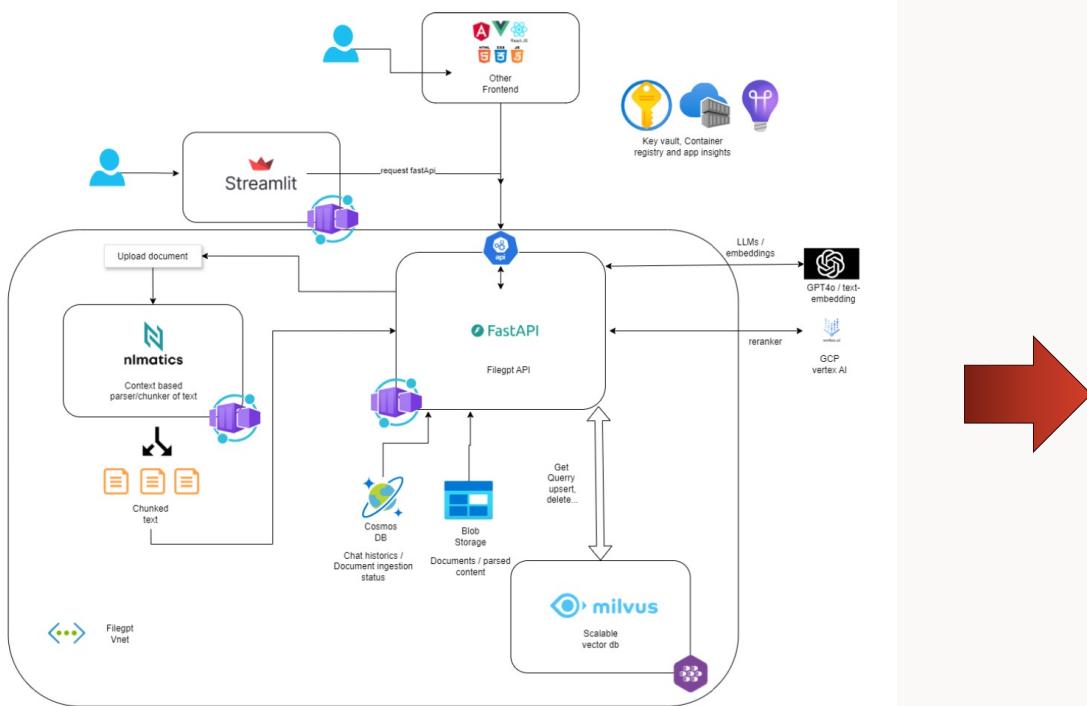


# A real architecture as example...

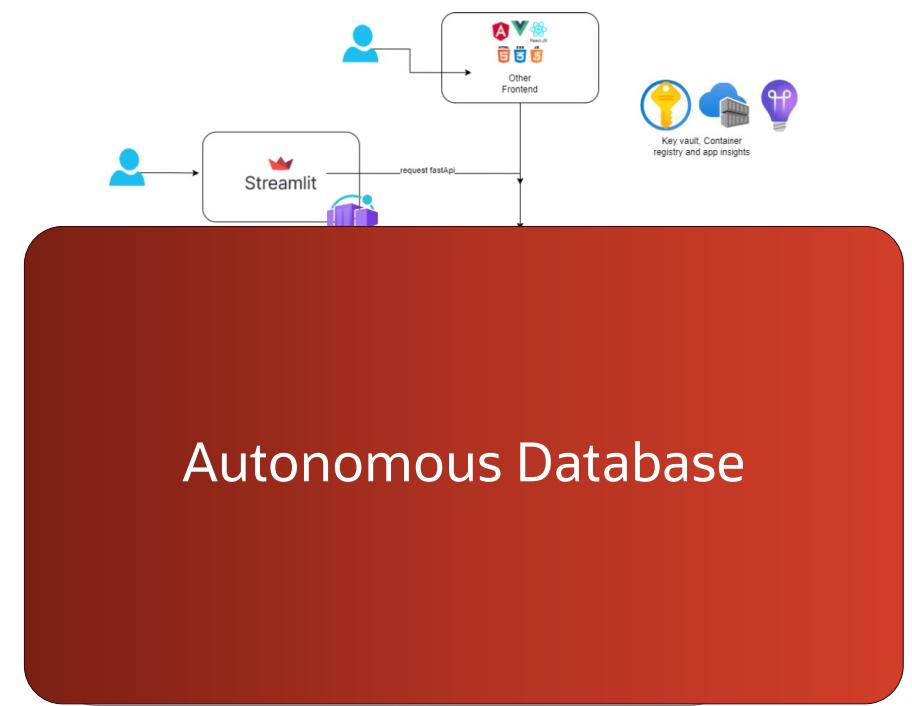


# A real architecture as example...

System Architecture



System Architecture

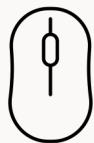


# The “Three Ways of RAG” using Autonomous AI Database



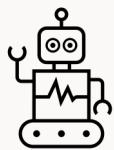
## **Write your own server code**

this is the recommended way when you want total control



## **No code / use embedded capabilities**

this is the recommended way when you are focused on getting a result fast



## **Agentic approach**

this is the recommended way when RAG is to be one part in a larger set of operations



# The “Three Ways of RAG” using Autonomous AI Database

*Write your own server code – project assumptions*

- Domain Knowledge stored as documents in an object store bucket
- Don't want/need to load the documents into the DB first (duplication/cost)
- No front-end dependencies (once I get the answer how it gets visualized is independent from it)



# The “Three Ways of RAG” using Autonomous AI Database

*Write your own server code – technology choices*

- Select AI
  - Essentially just two calls
  - Limited control on process settings
  - SQL interface
- Vector Index
  - Definitely more than two calls
  - Full control on chunking, embedding, results filtering, and answer creation
- Hybrid Vector Index
  - Complex but very powerful query syntax
  - Less control on the settings
  - Huge optimization (avoid expensive reranking!) through union of Text and Vector searches



# The “Three Ways of RAG” using Autonomous AI Database

No code approach – just load your documents

The screenshot shows the Oracle Applied AI Studio interface. On the left is a sidebar with navigation links: STUDIO (Getting Started, Template Gallery, My Custom Workflows), PRE-BUILT AGENTS (Knowledge Agent, Data Analysis Agent), UTILITIES (Agent Builder, Prompt Lab, Prompt2Dashboard, SQL Viewer, Datasets), SETTINGS (Data Sources, SSO, Users). The main area is titled "Capacity\_Management\_Agent" and describes it as a knowledge agent trained on capacity management documentation. It shows a message history between the user ("echelon\_admin@oracle.com (You)") and the "Knowledge Agent". The user asks, "How do I terminate an OCI tenancy request?", and the agent responds with a detailed answer. Below the response, there are links to "Requesting a New Commercial Subscription for Your Tenancy", "Terminating an Internal Account", and "Capacity or Tenancy Requests". A footer at the bottom says "How can I help you?" and has a "Feedback" button.

ORACLE® Applied AI Studio

echelon\_admin@oracle.com

STUDIO

- Getting Started
- Template Gallery
- My Custom Workflows

PRE-BUILT AGENTS

- Knowledge Agent
- Data Analysis Agent

UTILITIES

- Agent Builder
- Prompt Lab
- Prompt2Dashboard >
- SQL Viewer
- Datasets

SETTINGS

- Data Sources
- SSO
- Users

Capacity\_Management\_Agent

A knowledge agent trained on capacity management documentation.

Messages Chat History Help + New Chat

echelon\_admin@oracle.com (You) 7/29/2025, 12:01:10 PM

How do I terminate an OCI tenancy request?

Knowledge Agent 7/29/2025, 12:01:17 PM

To terminate an OCI tenancy request, you can follow these steps: Ensure that you have the necessary access permissions to manage OCI tenancies. Only users with the appropriate privileges, such as a tenancy administrator, can terminate tenancy requests. Verify that the tenancy request you want to terminate is in a state that can be terminated. Check the tenancy request status to ensure that it is not in a pending or provisioning state. Open a web browser and navigate to the OCI console. Sign in to the OCI console using your credentials. Navigate to the Tenancies page and locate the tenancy request you want to terminate. Click on the three vertical dots next to the tenancy request and select "Terminate" from the dropdown menu. Confirm that you want to terminate the tenancy request by clicking "Terminate" in the confirmation dialog box. Verify that the tenancy request has been terminated successfully. You can check the tenancy request status to ensure that it has been updated to "Terminated". It is essential to note that terminating a tenancy request will release the capacity associated with it, making it available for reclamation and reuse.

Sources

[Requesting a New Commercial Subscription for Your Tenancy](#)

[Terminating an Internal Account](#)

[Capacity or Tenancy Requests](#)

How can I help you?

# The “Three Ways of RAG” using Autonomous AI Database

## Agentic RAG – build your own RAG agent

The screenshot shows the Oracle Applied AI Studio interface. On the left, the sidebar includes sections for STUDIO (Getting Started, Template Gallery, My Custom Workflows), PRE-BUILT AGENTS (Knowledge Agent, Data Analysis Agent), UTILITIES (Agent Builder, Prompt Lab, Prompt2Dashboard, SQL Viewer, Datasets), and SETTINGS (Data Sources, SSO, Users, LLM Management). The main workspace displays a workflow titled "My Workflows / Diagnostics and Documentation Agent via...". The workflow consists of several components connected by arrows:

- A "MCP Server" component with a URL field containing "http://phoenix726584.dev3sub3.c".
- An "OCI Agent" component with a configuration field set to "oci-llama3.3".
- A "Tools" component.
- A "Custom Instructions" component with text: "Help users triage software bugs using natural language. You're a helpful assistant in a software team. You understand how to interact with the bug triage system to check bug status". It includes a "Save Prompt" button.
- A "Chat Input" component labeled "User message send in chat".
- A "Message" component.
- A "Prompt" component with a "Receiving Input" field and a "Temperature" slider set to 0.01.
- A "Chat Output" component labeled "Send message to chat".
- A "Message" component.

Annotations with blue circles highlight specific components: one on the "Tools" component, one on the "Custom Instructions" component, and one on the "Message" component at the bottom right.

## Recap & Take Aways

1

**Oracle Autonomous AI Database helps reduce complexity**

It is multicloud, secure and scalable

2

**Multiple technology paths**

Freedom of choice between code, no-code, and agentic approach

3

**Focus on the business case of the RAG solution**

Technology makes it simple, but be very clear on what business problem you are addressing



RAG in a Box  
Blog Article



Select AI



Vector Search



Hybrid Vector Search



**Your feedback is  
important.**

**Scan this QR Code or use  
the Mobile App to share  
your thoughts on this  
session.**



To contact us

---

**Massimo Castelli**



**Doug Hood**



# Thank you!

---



**ORACLE**

O